

Isocolon, parison, and comparative construal: A corpus test of formal balance

Brett Reynolds 

Humber Polytechnic & University of Toronto

30th May 2026

Abstract

Rhetorical theory often treats balanced form as more than ornament. A balanced pair of cola can present two units as candidates for comparison. A reader-side causal claim would need experimental evidence, but a production-side textual claim can be tested in corpus data. This paper tests whether adjacent discourse-relation pairs in the Georgetown University Multilayer corpus with enhanced Rhetorical Structure Theory annotations (GUM/eRST) show more **FORMAL BALANCE** when they construe two units as comparable co-members of a set. The analysis separates **ISOCOLON**, measured as length balance, from **PARISON**, measured as syntactic parallelism, and **LEXICAL ECHO**. In 17,456 adjacent relation pairs from 299 documents across 24 genres, a broad comparison and co-ordination family shows a positive adjusted increase in composite formal balance. The largest broad-family effects come from lists and disjunctions, where co-ordination is a strong grammatical null; the cleaner non-enumerative signal comes from adversative contrast.

The composite increase is mostly parison-like. Mean-only length controls also show length-balance effects, but those effects are not the main source of the composite signal. A matched co-ordination check shrinks the list/disjunction effect; an adversative-only check leaves contrast positive. Antithesis alone is near zero, and concession is negative. The result supports a restrained textual claim: formal balance is part of the profile of some comparable-set construals and varies across relations that could be called contrastive.

1 INTRODUCTION

Rhetorical theory has long treated figures as patterned forms with argumentative force. In that tradition, **ISOCOLON** joins two kinds of evidence that are easy to confuse: formal balance between adjacent units and a reader's sense that the units invite comparison. The motivating claim is simple. When writers place two clauses or phrases in balanced form, they may present the comparison between them as more explicit, pointed, or forceful.

This paper asks how much of that claim can be tested without a reader experiment. The answer depends on the explanatory level. At the reader-response level, the claim is causal: isocolon changes

*Contact: brett.reynolds@humber.ca

what readers notice or how forcefully they construe a relation. Corpus evidence alone can't show that. At the textual level, the claim is distributional: writers may preferentially use balanced form where they construe two units as comparable. That claim is testable.

The corpus test uses a narrower hypothesis. When a relation places two adjacent units in a shared comparison frame, the pair should show more formal balance than pairs expressing cause, condition, elaboration, sequence, or other non-comparative relations. The study uses the Georgetown University Multilayer (GUM) corpus and enhanced Rhetorical Structure Theory (eRST) annotations because they provide genre-diverse texts, discourse-relation labels, and signal annotation (Zeldes, 2017; Zeldes et al., 2025). The present analysis gives a first distributional test using GUM/eRST. The Rhetorical Structure Theory (RST) Signalling Corpus will allow later triangulation against a broader inventory of coherence-relation signals.

The pattern refines the rhetorical claim. Broadly comparative and co-ordinative relations are more formally balanced than matched non-target relations, but the largest broad-family effects occur exactly where a grammar-driven co-ordination account is strongest: lists and disjunctions. Ordinary adversative contrast is smaller but cleaner, because it survives inside an adversative-only comparison. Antithesis alone is near zero, and concession is negative. The pattern suggests that formal balance, especially parison-like syntactic parallelism, tracks some acts of comparative construal better than it tracks the inherited relation name ANTITHESIS.

2 RHETORICAL FIGURES AS FORM-FUNCTION PAIRINGS

The theoretical starting point is Fahnestock's account of rhetorical figures as form-function pairings. In her work on scientific argument and rhetorical style, figures are recurrent formal arrangements that can epitomize lines of reasoning (Fahnestock, 1999, 2011). That view gives the corpus prediction. If a figure is tied to an argumentative operation, then the figure should be distributed unevenly across the discourse contexts where that operation is being performed.

Fahnestock's discussion of the figures of argument is especially direct for this project. ISOCOLON concerns approximate equality in the length of adjacent cola, while PARISON concerns similarity of grammatical structure. The two often co-occur, but the properties differ. A pair can be close in length without being syntactically parallel, and it can be syntactically parallel without being closely balanced in length. The classical tradition links both to comparison and contrast because similar form can make two units cohere as an argumentative set (Fahnestock, 2011).

A figure can be a resource for making a relation salient without being obligatory whenever the relation is present. For the corpus test, uneven distribution is expected, and frequency by itself gives no direct measure of the rhetorical value of the figure.

The same point motivates computational work on rhetorical figures. Harris and Di Marco (2017) argue that figural form is often iconically related to discourse function, so surface patterning can provide evidence about pragmatic and argumentative actions. Harris et al. (2018) develop an annotation scheme for figures in part because figure identification requires attention to both formal regularity and rhetorical role. Lawrence et al. (2017) make the parallel point for argument mining: rhetorical figures can identify regions where argumentative work is being done, but only if formal detection is connected to the content and relation being advanced.

Together, these studies support the central assumption of the present study. The relevant object is the co-occurrence of surface balance with a discourse relation that makes two units available for

comparison. The operational question is whether formal balance, separated into isocolonic length balance, parison-like syntactic parallelism, and lexical echo, is higher in the relations where rhetorical theory predicts that balance has value.

3 FROM CANDIDATE FORM TO RHETORICAL FUNCTION

The primary methodological challenge is over-attribution. Repetition and parallelism are common in ordinary text, and only some instances have the rhetorical force that figure theory cares about. Work on automatic chiasmus detection makes the problem vivid. Dubremetz and Nivre (2018) show that trivial repetition patterns are easy to find, but repetitions with rhetorical effect are much harder to separate from accidental or irrelevant patterns. They treat figure detection as a ranking problem instead of a clean binary classification problem.

That lesson carries over to isocolon and parison. Exact repetition, list formatting, dialogue formulas, quoted slogans, and segmentation artifacts can all look formally balanced. Some are rhetorically relevant, but many are not instances of the target phenomenon. Recent survey work makes the broader point: rhetorical figure detection remains limited by small datasets, inconsistent definitions, and over-reliance on rule-based surface cues (Kühn & Mitrović, 2024; Kühn et al., 2024). A credible corpus test should use surface scores cautiously, validate them against human judgment, and avoid treating high balance as equivalent to a gold figure label.

The analysis below follows that strategy. The composite formal-balance score is a continuous outcome in the distributional analysis and a ranking score in validation. It ranks adjacent relation pairs by how much balance they display. The qualitative audit then asks whether high-scoring examples are plausible rhetorical-balance candidates after excluding exact repetition, dialogue formulas, and other false-positive types. This more modest target than automatic figure recognition is the right target for a first rhetorical-function test.

4 CORPUS AND MEASUREMENT

The corpus is the English Georgetown University Multilayer corpus with enhanced Rhetorical Structure Theory annotations (GUM/eRST). GUM is a genre-diverse corpus with richly annotated English texts (Zeldes, 2017). eRST extends RST-style discourse annotation with graph structures and explicit signal annotation across spoken and written genres (Zeldes et al., 2025). The project needs discourse relations as well as adjacent sentences. It also needs enough genre variation to separate a rhetorical effect from a register-specific habit.

The unit of analysis is an adjacent discourse-relation pair. The dataset contains 17,456 such pairs from 299 documents across 24 genres. These discourse units serve as candidate cola-like spans, with no claim that they're classical cola in a strict prosodic sense.

A relation enters the analysis if eRST marks it as explicit or implicit, both spans have token intervals, and the two intervals are adjacent in document order. Non-adjacent linked spans are excluded. Concurrent graph relations are kept as pairwise relation rows when their spans meet this adjacency criterion.

The DISRPT relation export used here omits the full eRST graph-node inventory from the derived table. The diagnostics in Table 1 are interval-level checks instead of a complete graph-theoretic audit. The extracted set contains 16,278 rows with collection label GUM and 1,178 with collection label *GENTLE*, the out-of-domain challenge set; the 24 genres are the document-genre labels in the

extracted files.

Extraction diagnostic	Count
Valid explicit/implicit relation rows	34,976
Adjacent relation rows used in analysis	17,456
Non-adjacent relation rows excluded	17,520
gapped rows excluded	15,349
interleaved rows excluded	2,171
Adjacent rows with discontinuous token intervals	1,963
Exact duplicate adjacent span pairs	0
Adjacent explicit rows	4,943
Adjacent implicit rows	12,513

Table 1: Extraction diagnostics for the adjacent-pair analysis. Counts come from token intervals in the DISRPT relation export. Exact duplicate span pairs are counted by document, unit-1 tokens, and unit-2 tokens.

For each pair, I computed three surface measures. The first is *ISOCOLON*, operationalized as length balance over token, character, and estimated syllable length. The second is *PARISON*, operationalized as similarity in dependency and part-of-speech profiles. The third is *LEXICAL ECHO*, based on repeated or closely corresponding lexical material. A composite score combines the three as a graded measure of formal balance. The score is evidence about balance, not a figure label for the pair.

The component scoring is deliberately simple. For any two units with positive length, a balance component is $1 - |n_1 - n_2| / \max(n_1, n_2)$. The length score averages this quantity over word tokens, characters, and estimated syllables. The syntax score averages normalized edit similarity over universal part-of-speech (UPOS), language-specific part-of-speech (XPOS), and dependency-relation sequences read in token order. I treat this as a surface syntactic profile, not a tree-alignment measure. The lexical score averages lemma-sequence edit similarity and Jaccard overlap over content lemmas.

The GUM/eRST files use the Conference on Natural Language Learning format for Universal Dependencies (CoNLL-U), which records tokenization, lemmas, part-of-speech tags, and dependency labels. Punctuation tokens are excluded from the word, syntax, and lexical sequences, while colon and semicolon position are kept as separate controls. Syllable counts are approximate vowel-group counts.

The composite is

$$0.40 \times \text{length} + 0.45 \times \text{syntax} + 0.15 \times \text{lexical}.$$

These weights define only the composite formal-balance score. The components are also reported separately, so the analysis can distinguish isocolonic length balance from parison-like syntactic parallelism.

The target relations were defined before the final robustness checks as analysis groups, separate from a new annotation layer. Table 2 gives the mapping. The broad target family includes *adversative-antithesis*, *adversative-contrast*, *joint-disjunction*, and *joint-list*. This family tests whether formal balance rises in relations that place units in a comparable or co-ordinated frame. The narrow target family includes only *adversative-antithesis* and *adversative-contrast*. Individual-relation analyses separate the broad effect into its component labels.

Analysis role	Labels, rationale, and caveat
Broad target	<i>adversative-antithesis; adversative-contrast; joint-disjunction; joint-list</i> . Contrast, alternatives, and co-enumeration place units into a shared frame. The joint labels are co-ordination-rich and also the strongest grammar-driven null.
Narrow adversative target	<i>adversative-antithesis; adversative-contrast</i> . Tests whether the broad result survives without list and disjunction labels.
Cleaner non-enumerative test	<i>adversative-contrast</i> . Tests contrast against other adversative relations, where list syntax is not the main explanation.
Contrastive control	<i>adversative-concession</i> . Adversative but often asymmetric; predicted to be weaker or negative.
Surface positive control	<i>restatement-repetition</i> . Expected to score high by form alone; used as a false-positive diagnostic, separate from the comparison target.

Table 2: Mapping from eRST relation labels to analysis roles. The table defines the operational grouping. Independent annotation of comparable-set construal remains future work.

The primary comparisons adjust for the most obvious structural confounds. Where possible, pairs are compared within documents. The models also include controls for explicit signalling, sentence adjacency, punctuation, mean length, and maximum length. Those controls are needed for the composite score because longer spans have more opportunities for structural mismatch. For the length subscore, the pair of length controls overcontrols the target quantity: token-length balance is $2\mu/m - 1$, where μ is mean span length and m is maximum span length. The isocolon subscore is reported with mean-length control and without maximum-length control.

A genre-varying analysis then checks whether the effect is concentrated in one or two genres or survives partial pooling across genres.

The adjusted estimates are ordinary least-squares target coefficients computed by residualizing both the target indicator and the outcome against the controls. The controls are an intercept, explicit relation status, same-sentence adjacency, colon punctuation, semicolon punctuation, standardized log mean length, standardized log maximum length, and document fixed effects. The genre-varying analysis first fits the same adjusted effect within each genre, then uses a normal empirical-Bayes model with genre effects distributed around a shared mean and between-genre standard deviation.

Robustness checks add uncertainty estimates clustered by document, mean-only length-balance models, a joint-only comparison, an adversative-only comparison, and score-weight sensitivity checks.

5 VALIDATION

The composite score ranks candidates for audit. Against the small number of eRST *syn-prl* signal labels in the dataset, the composite score has an area under the receiver-operating characteristic curve (AUC) of 0.731, with an approximate Hanley–McNeil interval from 0.664 to 0.797. The parison subscore is slightly higher at 0.754 [0.689, 0.819], lexical echo is 0.719 [0.652, 0.786], and isocolonic length balance is 0.646 [0.577, 0.715] (Table 3).

At the 95th percentile, precision is only 3.1%, but recall is 37.0%, a lift of 7.4 over prevalence. At the 97.5th percentile, precision rises to 5.5%, recall is 32.9%, and lift is 13.1. Those numbers are poor for automatic labelling but adequate for ranking rare candidates for audit. The validation is also narrow: the *syn-prl* labels validate a surface parallelism signal. They don’t validate rhetorical function itself.

Score	AUC	Gold mean	Non-gold mean
Formal balance	0.731 [0.664, 0.797]	0.502	0.327
Isocolon	0.646 [0.577, 0.715]	0.683	0.560
Parison	0.754 [0.689, 0.819]	0.441	0.209
Lexical echo	0.719 [0.652, 0.786]	0.202	0.057

Table 3: Score validation against the 73 eRST *syn-prl* signal labels in 17,456 adjacent relation pairs. Brackets give approximate Hanley–McNeil 95% intervals for AUC.

The qualitative audit carries the rhetorical-function check. Eighty candidate pairs were coded by the author in a single-author, non-blind pilot audit: row strata and scores were visible during coding. Automated reviews flagged borderline cases for reinspection after coding. The audit yielded 33 yes, 43 no, and 4 uncertain judgments. I use it to identify false-positive classes and illustrative examples, with no claim of independent interrater reliability.

High-scoring broad-target candidates performed well: the first batch had 8 yes judgments out of 10 high broad-but-not-narrow items, and the second had 6 yes judgments out of 6 high broad-target items. By contrast, the restatement/repetition controls had 0 yes judgments out of 9. Exact repetition is the strongest surface correlate of formal balance and a separate false-positive class for this study.

The audit also clarified the boundary of the category. Bare identity, apposition, dialogue repair, and exact repetition can be rhetorically relevant apart from isocolon. Conversely, weak length balance can still matter when the two units share a frame strongly enough to invite slot substitution. The coding decisions treat isocolon as a component of a broader rhetorical-balance judgment, with false-positive classes instead of a mechanically recoverable surface pattern.

Sample stratum	<i>n</i>	Yes	No	Unc.	Main issue
High broad target	16	14	2	0	Slot substitution
Mid broad target	12	7	4	1	Weak frames
Narrow adversative	14	6	8	0	Label mismatch
High non-target/artifact	19	4	12	3	Dialogue/apposition
Concession candidates	5	1	4	0	Asymmetry
Repetition controls	9	0	9	0	Exact identity
Baseline non-target	5	1	4	0	Weak or no frame

Table 4: Human qualitative audit summary for 80 candidate pairs.

Type	Unit 1	Unit 2	Judgment
Slot substitution	<i>You can use this on twist outs.</i>	<i>You can use this on braid outs.</i>	Yes
Adversative contrast	<i>so that the broad sides of the cuttings face east and west</i>	<i>while the slim sides face north and south.</i>	Yes
Repetition	<i>Get your feet off of there.</i>	<i>Get your feet off of there.</i>	No
Label mismatch	<i>We are not setting out to turn them into programmers or data crunchers,</i>	<i>but want to boost their knowledge level</i>	No

Table 5: Illustrative audit examples. The final row is an *adversative-antithesis* relation label. It lacks the balanced opposed cola of classical antithesis.

6 RESULTS

The broad target family shows a positive adjusted increase in composite formal balance. Across all adjacent pairs, the adjusted effect is 0.0415 with a document-clustered interval from 0.0356 to 0.0474. That's about 0.28 standard deviations of the observed composite score. The subscales show that the composite effect is driven most strongly by parison: the adjusted syntax effect is 0.0770, compared with 0.0244 for lexical echo. With mean-length control only, the broad isocolon effect is 0.0664. The component split identifies which part of balanced form carries each signal.

Table 6 reports adjusted target coefficients for the composite and the three components separately.

The broad-family coefficient is an omnibus check. Its strongest labels are also the labels where co-ordination makes parallel form easiest, so the cleaner rhetorical test comes from narrower comparisons.

The narrow adversative family is smaller. For *adversative-antithesis* plus *adversative-contrast*, the adjusted composite effect is 0.0130, with a confidence interval from 0.0048 to 0.0212.

The individual labels explain the difference. *joint-disjunction* has an adjusted formal-balance effect of 0.0482, *joint-list* has an adjusted effect of 0.0439, and *adversative-contrast* has the smaller but more diagnostic effect of 0.0241. *adversative-antithesis* is near zero at 0.0039, with an interval that crosses zero. I read this as a mismatch between a discourse-relation label and the classical figure. It gives no evidence against antithesis as a rhetorical resource. *adversative-concession* is negative at -0.0131.

The relation-level scan also identifies an important positive control. *restatement-repetition* has the strongest adjusted composite effect, 0.1902, with especially large lexical and parison effects. That is exactly what a surface score should find: exact identity is highly balanced, with lexical echo stronger than the syntactic effect. The human audit shows why it needs separate treatment. Repetition can signal clarification, identity, repair, emphasis, or other rhetorical work; exact repetition remains distinct from balanced comparison. The analysis treats restatement/repetition as a control and diagnostic, separate from the comparison hypothesis.

Genre-varying estimates point in the same direction. The broad target has a partially pooled mean effect of 0.0390 across 20 genres. The *joint-list* estimate is 0.0398 across 17 genres, and *adversative-contrast* is 0.0291 across the genres where there is enough evidence to estimate it. The narrow adversative estimate is 0.0118. The estimates vary across genres, and the broad effect appears in multiple genres.

Target	<i>n</i>	Formal	Isocolon	Parison	Lexical
Broad family	2,462	0.041	0.066	0.077	0.024
Narrow adversative	605	0.013	0.047	0.024	0.017
<i>joint-list</i>	1,690	0.044	0.065	0.082	0.020
<i>joint-disjunction</i>	167	0.048	0.036	0.086	0.040
<i>adversative-contrast</i>	262	0.024	0.073	0.045	0.028
<i>adversative-antithesis</i>	343	0.004	0.025	0.008	0.009
<i>adversative-concession</i>	679	-0.013	0.011	-0.027	-0.007
<i>restatement-repetition</i>	321	0.190	0.170	0.273	0.325

Table 6: Selected adjusted target coefficients on composite formal balance and its components. The *n* column gives the number of target rows. ISOCOLON is length balance, PARISON is syntactic parallelism, and LEXICAL ECHO is repeated or corresponding lexical material. Formal, parison, and lexical estimates use document fixed effects and controls for explicitness, sentence adjacency, punctuation, mean length, and maximum length. Isocolon estimates omit maximum length because mean and maximum length nearly determine the length-balance score. The restatement row is a positive control, separate from the comparison target.




Figure 1: Adjusted target coefficients for FORMAL BALANCE, ISOCOLON, PARISON, and LEXICAL ECHO. Points are estimates; horizontal intervals are document-clustered 95% intervals. The ISOCOLON estimates use mean-length controls and omit maximum length; the other estimates use the main control specification.

6.1 ROBUSTNESS CHECKS

The primary formal-balance results survive document-clustered uncertainty estimates. With document clustering, the broad-family effect is 0.041 with an interval from 0.036 to 0.047. The *joint-list* effect is 0.044 [0.037, 0.051], *joint-disjunction* is 0.048 [0.030, 0.067], and *adversative-contrast* is 0.024 [0.013, 0.035]. *adversative-antithesis* remains near zero, 0.004 [-0.007, 0.015].

The length-balance result uses a different control specification. Mean span length and maximum span length nearly determine the length-balance score, so using both controls suppresses the quantity being estimated. With mean-length control but without maximum-length control, the broad-family effect on length balance is 0.066 [0.056, 0.077], and *adversative-contrast* is 0.073 [0.050, 0.096]. The corresponding no-length-controls estimates are nearly identical: 0.067 for the broad family and 0.071 for *adversative-contrast*. The length-balance claim is reported under mean-length control.

The co-ordination alternative explains part of the pattern. Within *joint-** relations only, *joint-list* plus *joint-disjunction* has a smaller formal-balance effect of 0.017 [0.006, 0.028] against other joint relations. The corresponding parison effect is 0.029 [0.010, 0.048], while the isocolon effect is only 0.003 [-0.007, 0.013].

Within adversative relations only, *adversative-contrast* remains positive at 0.030 [0.014, 0.047], again led by parison. These checks give the co-ordination explanation a direct test. The joint-only baseline consists of *joint-sequence* and *joint-other*; the baseline remains multinuclear and partly co-ordinative. The adversative-only contrast is cleaner. It compares *adversative-contrast* against other adversative relations.

The stratified permutation check shuffled target labels 500 times within strata defined by document, explicitness, same-sentence status, colon punctuation, semicolon punctuation, and quartile of mean word length. The model specification was otherwise the same as in the main composite analysis. This partial null leaves local phrase type and dependency-root category unmatched. It tests whether the result follows only from the allocation of target labels across documents, lengths, punctuation, and explicitness. The observed broad-family effect, the list/disjunction effect, and the adversative-contrast effect all fall above the 97.5th percentile of their shuffled null distributions (Table 8).

The composite result is stable across the hand-set weights. Across 66 grid-based nonnegative weightings of length, syntax, and lexical echo, the broad-family effect is positive with a positive interval in every specification. For *adversative-contrast*, 65 of 66 grid specifications are positive and 59 have positive intervals.

Check	Target/outcome	Estimate	Clustered 95% interval
Main	Broad formal balance	0.041	[0.036, 0.047]
Main	<i>joint-list</i> formal balance	0.044	[0.037, 0.051]
Main	<i>adversative-contrast</i> formal balance	0.024	[0.013, 0.035]
Mean-length only	Broad isocolon	0.066	[0.056, 0.077]
Mean-length only	<i>adversative-contrast</i> isocolon	0.073	[0.050, 0.096]
Joint-only	List/disjunction formal balance	0.017	[0.006, 0.028]
Joint-only	List/disjunction parison	0.029	[0.010, 0.048]
Adversative-only	Contrast formal balance	0.030	[0.014, 0.047]
Adversative-only	Contrast parison	0.059	[0.030, 0.088]

Table 7: Selected robustness checks. The joint-only rows compare *joint-list* and *joint-disjunction* against other *joint*-* relations. Mean-length-only rows control log mean length and omit log maximum length. The adversative-only rows compare *adversative-contrast* against other adversative relations.




Figure 2: Targeted robustness checks and stratified permutation nulls. Panel A shows adjusted target coefficients with document-clustered 95% intervals. Panel B compares observed composite effects with the null mean and central 95% interval from stratified target-label shuffles.

Target	Observed	Null mean	Null 95% interval	<i>p</i>
Broad family	0.041	0.012	[0.008, 0.016]	0.002
List/disjunction	0.046	0.013	[0.010, 0.017]	0.002
<i>adversative-contrast</i>	0.024	0.008	[-0.001, 0.017]	0.002

Table 8: Stratified permutation nulls for the composite formal-balance effect. Target labels are shuffled within document, explicitness, same-sentence status, colon/semicolon status, and mean-length quartile strata. The *p* values are two-sided empirical values from 500 permutations, with the standard plus-one correction.

7 INTERPRETATION

The findings support a restricted version of the rhetorical-function claim. Balanced form is best treated as a selective way of presenting units as comparable. It's optional even when comparison is present. That view predicts an uneven distribution: balance should cluster where matched presentation has rhetorical value, while remaining absent in many ordinary cases. The corpus result fits that version, but unevenly. Lists and disjunctions supply the largest broad-family effects; adversative contrast supplies the cleaner non-enumerative case.

Here COMPARABLE SET is a descriptive cover term for co-membership relations, including bare co-enumeration. The term allows list structure without treating every list as comparison for its own sake. The cleanest non-enumerative signal in the analysis is *adversative-contrast*.

The component results then refine the claim. The dominant effect is parison-like more than isocolonic in the narrow length-balance sense. The evidence is weaker for labels that look rhetorically relevant in the taxonomy but don't necessarily require balanced presentation in the text. In particular, *adversative-antithesis* as an eRST relation label only partly overlaps with antithesis in the classical theory of figures. The conservative conclusion leaves antithesis itself intact: the corpus label *adversative-antithesis* is a poor direct proxy for the classical figure.

That point remains a limitation as well as a finding. The 80-pair audit contains only eight rows with the *adversative-antithesis* label. These rows support the label-mismatch interpretation, but they're too few to separate label mismatch from a second possibility: genuine semantic antitheses may be expressed through opposition instead of lexical overlap or surface parallelism. A targeted follow-up should mark classical antithesis inside the eRST *adversative-antithesis* set and then test that subset directly.

The composite formal-balance score is associated with a broader act of placing units into a shared frame. That frame can be contrastive, but it can also be co-ordinative, disjunctive, or list-like. In the latter cases, co-ordination itself favours matched daughters and weakens the rhetorical interpretation. Conversely, some adversative relations have little balance because their function is concession, qualification, or discourse management instead of matched presentation.

Corpus data record ordinary production, not an ideal of trained rhetorical practice. The present evidence shows where balance appears in this corpus. It leaves open where balanced form would improve a text.

The co-ordination caveat is substantial. Lists and disjunctions are places where parallel syntax is grammatically cheap, so the broad effect cannot be read as a pure rhetorical choice. The joint-only robustness check shrinks the effect, which is what that objection predicts. The effect remains positive, and *adversative-contrast* remains positive inside the adversative family and under a stratified

permutation null. The safer interpretation is that comparable-set relations attract formal balance beyond genre, document, length, and punctuation differences, while much of that balance follows the syntax of co-ordination.

The pattern also clarifies the scope of the corpus study. The evidence is textual and distributional. It concerns where balanced form occurs, not how readers experience those pairs. A high score by itself still falls short of a rhetorical-figure label. The analysis shows that a graded measure of formal balance is systematically higher in particular discourse-relation contexts. The human audit illustrates why many high-scoring cases are rhetorically relevant and why this remains pilot evidence, pending an independently validated annotation layer.

The RST Signalling Corpus will add a later triangulation point once the data is available. Work on discourse signalling has shown that coherence relations are often signalled by more than discourse markers, including lexical, syntactic, graphical, semantic, and genre-based cues. That point is developed in the original signalling annotation work and in later studies of RST signals and multiple signals (Das & Taboada, 2018a, 2018b, 2019; Taboada & Das, 2013). This literature gives the project a stronger comparative baseline. If formal balance behaves like a discourse signal, then it should be testable alongside other signal types as well as an isolated figure score.

8 CONCLUSION

A corpus study cannot prove that isocolon makes readers perceive comparison as more forceful, and it leaves rhetorical pedagogy untested. It can test a cleaner textual claim: writers use formal balance unevenly across discourse relations, and the relations most associated with comparable-set construals should show more of it. The present results support that claim for a broad family of comparison-like and co-ordinative relations, with the largest effects in lists and disjunctions and the cleaner non-enumerative effect in adversative contrast.

The analysis gives selective support to the inherited figure label *ANTITHESIS*. The narrow adversative effect is small, and antithesis as a discourse-relation label is close to zero. That finding sharpens the mapping between relation labels and classical figures. The strongest signal is parison-like parallelism; isocolon narrowly understood as length balance contributes less to the main composite effect. Balanced form remains a rhetorical resource, and the present corpus evidence supports a claim about comparable-set construals and formal balance more securely than a claim about equal cola alone.

The robustness checks sharpen that conclusion. Length balance appears under mean-length conditioning without maximum length, while the main composite model is led by parison. The joint-only comparison also shows that a large share of the broad effect reflects co-ordination itself. The stratified permutation check places the observed effects above document, length, punctuation, and explicitness-matched nulls. What remains is a modest, distributional claim: comparable-set relations prefer matched form in ordinary corpus production, and the clearest measurable form in this corpus is syntactic parallelism.

DATA AND CODE AVAILABILITY

The analysis scripts, coding materials, and derived audit summaries are kept in the project repository: <https://github.com/BrettRey/isocolon-and-comparative-construal>. The raw GUM/eRST source files are excluded from the repository; the scripts expect a local corpus checkout under `data/raw/gum-erst`. RST Signalling Corpus data will be added after access is available.

ACKNOWLEDGEMENTS

This paper was drafted with the assistance of large language models. All content has been reviewed and revised by the author, who takes full responsibility for the final text.

REFERENCES

- Das, D., & Taboada, M. (2018a). RST signalling corpus: A corpus of signals of coherence relations. *Language Resources and Evaluation*, 52(1), 149–184. <https://doi.org/10.1007/s10579-017-9383-x>
- Das, D., & Taboada, M. (2018b). Signalling of coherence relations in discourse, beyond discourse markers. *Discourse Processes*, 55(8), 743–770. <https://doi.org/10.1080/0163853X.2017.1379327>
- Das, D., & Taboada, M. (2019). Multiple signals of coherence relations. *Discours*, 24. <https://doi.org/10.4000/discours.10032>
- Dubremetz, M., & Nivre, J. (2018). Rhetorical figure detection: Chiasmus, epanaphora, epiphora. *Frontiers in Digital Humanities*, 5, 10. <https://doi.org/10.3389/fdigh.2018.00010>
- Fahnestock, J. (1999). *Rhetorical figures in science*. Oxford University Press. <https://doi.org/10.1093/os0/9780195117509.001.0001>
- Fahnestock, J. (2011). *Rhetorical style: The uses of language in persuasion*. Oxford University Press. <https://doi.org/10.1093/acprof:os0/9780199764129.001.0001>
- Harris, R. A., & Di Marco, C. (2017). Rhetorical figures, arguments, computation. *Argument & Computation*, 8(3), 211–231. <https://doi.org/10.3233/AAC-170030>
- Harris, R. A., Di Marco, C., Ruan, S., & O'Reilly, C. (2018). An annotation scheme for rhetorical figures. *Argument & Computation*, 9(2), 155–175. <https://doi.org/10.3233/AAC-180037>
- Kühn, R., & Mitrović, J. (2024). The Elephant in the Room: Ten challenges of computational detection of rhetorical figures. *Proceedings of the 4th Workshop on Figurative Language Processing (FigLang 2024)*, 45–52. <https://doi.org/10.18653/v1/2024.figlang-1.6>
- Kühn, R., Mitrović, J., & Granitzer, M. (2024). Computational approaches to the detection of lesser-known rhetorical figures: A systematic survey and research challenges. <https://doi.org/10.48550/arXiv.2406.16674>
- Lawrence, J., Visser, J., & Reed, C. (2017). Harnessing rhetorical figures for argument mining: A pilot study in relating figures of speech to argument structure. *Argument & Computation*, 8(3), 289–310. <https://doi.org/10.3233/AAC-170026>
- Taboada, M., & Das, D. (2013). Annotation upon annotation: Adding signalling information to a corpus of discourse relations. *Dialogue & Discourse*, 4(2), 249–281. <https://doi.org/10.5087/dad.2013.211>
- Zeldes, A. (2017). The GUM corpus: Creating multilayer resources in the classroom. *Language Resources and Evaluation*, 51(3), 581–612. <https://doi.org/10.1007/s10579-016-9343-x>
- Zeldes, A., Aoyama, T., Liu, Y. J., Peng, S., Das, D., & Gessler, L. (2025). eRST: A signaled graph theory of discourse relations and organization. *Computational Linguistics*, 51(1), 23–72. https://doi.org/10.1162/coli_a_00538